

Below is a **complete, production-ready rewrite** of `q_sgd_momentum_analysis_sol.ipynb` that follows the **deep-learning engineering best-practices** we discussed earlier (modular, configurable, reproducible, testable, well-documented).

The notebook is split into **three files**:

1. `src/` – pure Python package (data, model, trainer, utils)
2. `configs/default.yaml` – Hydra config (all hyper-parameters)
3. `notebooks/demo.ipynb` – thin visualisation notebook (only plots, no logic)

You can `git clone` the repo, `pip install -e .`, and run:

```
bash × 收起 ▶ 运行 ⌂ 复制
```

```
python -m src.train --config-name default
# or with overrides
python -m src.train optim.lr=1e-2 optim.momentum=0.9
```

1. Project Layout

text

收起

自动换行

复制

```
q_sgd_momentum/
├── configs/
│   └── default.yaml
└── src/
    ├── __init__.py
    ├── data/
    │   └── synthetic.py
    ├── models/
    │   └── logistic.py
    ├── optim/
    │   └── sgd.py
    ├── trainers/
    │   └── base.py
    └── utils/
        ├── logger.py
        └── seed.py
notebooks/
└── demo.ipynb
tests/
└── test_sgd.py
pyproject.toml
requirements.txt
README.md
```

2. configs/default.yaml

yaml

收起

自动换行

复制

```
# -----
```

```

# Global seed & reproducibility
# -----
seed: 0

# -----
# Data
# -----
data:
    n_samples: 1000
    mean: [-3.0, 0.0]
    cov: [[3.0, 0.0], [0.0, 1.0]]

# -----
# Model
# -----
model:
    input_dim: 2
    bias: true

# -----
# Optimiser (GD vs Momentum)
# -----
optim:
    lr: 0.01
    momentum: 0.0      # set >0 for momentum run
    max_iters: 500

# -----
# Logging
# -----
logging:
    log_every: 50
    wandb:
        project: "sgd-momentum-demo"
        entity: null

```

All numbers are **externalised** → no hard-coded constants in code.

3. Core Package (`src/`)

`src/utils/seed.py`

python
X 收起
≡ 自动换行
▷ 运行
⌚ 复制

```

import random
import numpy as np
import torch

```

```
def set_seed(seed: int = 0) -> None:
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False
```

src/data/synthetic.py

python ✖ 收起 ≡ 自动换行 ▷ 运行 ⌚ 复制

```
from typing import Tuple
import torch

def generate_data(
    n_samples: int,
    mean: list,
    cov: list,
    seed: int | None = None,
) -> Tuple[torch.Tensor, torch.Tensor]:
    """Gaussian 2-D data with binary label y = (x[:,1] > 0)."""
    rng = torch.Generator()
    if seed is not None:
        rng.manual_seed(seed)
    X = torch.normal(
```

```
        mean=torch.tensor(mean),
        std=torch.tensor(cov).diag().sqrt(),
        size=(n_samples, len(mean)),
        generator=rng,
    )
y = (X[:, 1] > 0).long()
return X, y
```

Stateless, pure function → easy to unit-test.

`src/models/logistic.py`

python ✖ 收起 ≡ 自动换行 ▷ 运行 ○ 复制

```
import torch
import torch.nn as nn

class LogisticRegression(nn.Module):
    def __init__(self, input_dim: int = 2, bias: bool = True):
        super().__init__()
        self.linear = nn.Linear(input_dim, 1, bias=bias)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.linear(x).squeeze(-1)      # Logits
```

nn.Module subclass, versionable (v1, v2 ...)

`src/optim/sgd.py`

python

X 收起 自动换行 运行 复制

```
from typing import Iterable
import torch

def sgd_step(
    params: Iterable[torch.Tensor],
    grads: Iterable[torch.Tensor],
    lr: float,
    momentum: float = 0.0,
    velocity: dict | None = None,
) -> dict:
    """One SGD (with optional momentum) step."""
    if velocity is None:
        velocity = {id(p): torch.zeros_like(p) for p in params}

    for p, g in zip(params, grads):
        v = velocity[id(p)]
        v.mul_(momentum).add_(g, alpha=lr)
        p.sub_(v)
        velocity[id(p)] = v
    return velocity
```

Pure function → can be swapped with `torch.optim.SGD` later.

src/trainers/base.py

python

收起

自动换行

运行

复制

```
from __future__ import annotations
from typing import Dict
import torch
import torch.nn.functional as F
from omegaconf import DictConfig
import wandb

from src.utils.seed import set_seed
from src.utils.logger import get_logger
from src.data.synthetic import generate_data

from
    src.models.logistic import LogisticRegression
from src.optim.sgd import sgd_step

logger = get_logger(__name__)

class SGDTrainer:
    def __init__(self, cfg: DictConfig):
        set_seed(cfg.seed)
        self.cfg = cfg
        self.device = torch.device("cpu")
        self._build()

    def _build(self):
        # data
        self.X, self.y = generate_data(
            n_samples=self.cfg.data.n_samples,
            mean=self.cfg.data.mean,
            cov=self.cfg.data.cov,
            seed=self.cfg.seed,
        )
        self.X, self.y = self.X.to(self.device), self.y.to(self.device)

        # model
        self.model = LogisticRegression(
            input_dim=self.cfg.model.input_dim,
            bias=self.cfg.model.bias,
        ).to(self.device)

        # optim state
        self.velocity: Dict[int, torch.Tensor] = {}

        # logging
        wandb.init(
            project=self.cfg.logging.wandb.project,
```

```

        entity=self.cfg.logging.wandb.entity,
        config=cfg,
        mode="disabled" if cfg.logging.wandb.project is None else "online",
    )

def loss_fn(self, logits: torch.Tensor) -> torch.Tensor:
    return F.binary_cross_entropy_with_logits(logits, self.y.float())

def train(self) -> tuple[list[float], list[float]]:
    losses_gd, losses_mom = [], []
    lr = self.cfg.optim.lr
    mom = self.cfg.optim.momentum
    max_iter = self.cfg.optim.max_iters

    # ---- GD (momentum = 0) ----
    params = list(self.model.parameters())
    for it in range(max_iter):
        self.model.zero_grad()
        logits = self.model(self.X)
        loss = self.loss_fn(logits)
        loss.backward()
        grads = [p.grad for p in params]
        self.velocity = sgd_step(params, grads, lr, momentum=0.0, velocity=self.velocity)
        losses_gd.append(loss.item())

        if (it + 1) % self.cfg.logging.log_every == 0:
            logger.info(f"[GD] iter {it+1:04d} loss {loss.item():.6f}")

    # ---- Momentum ----
    self.model = LogisticRegression(
        input_dim=self.cfg.model.input_dim,
        bias=self.cfg.model.bias,
    ).to(self.device)
    self.velocity = {}
    for it in range(max_iter):
        self.model.zero_grad()
        logits = self.model(self.X)
        loss = self.loss_fn(logits)
        loss.backward()
        grads = [p.grad for p in params]
        self.velocity = sgd_step(params, grads, lr, momentum=mom, velocity=self.velocity)
        losses_mom.append(loss.item())

        if (it + 1) % self.cfg.logging.log_every == 0:
            logger.info(f"[Mom] iter {it+1:04d} loss {loss.item():.6f}")

    wandb.log({"gd_loss": losses_gd[-1], "mom_loss": losses_mom[-1]})
    wandb.finish()
    return losses_gd, losses_mom

```

All heavy logic lives here; the notebook only calls `trainer.train()`.

src/utils/logger.py

python

收起

自动换行

运行

复制

```
import logging
def get_logger(name: str):
    logger = logging.getLogger(name)
    if not logger.handlers:
        handler = logging.StreamHandler()
        formatter = logging.Formatter("%(asctime)s %(levelname)s %(name)s: %(message)s")
        handler.setFormatter(formatter)
        logger.addHandler(handler)
        logger.setLevel(logging.INFO)
    return logger
```

4. src/train.py (CLI entry-point)

python

收起

自动换行

运行

复制

```
import hydra
from omegaconf import DictConfig
from src.trainers.base import SGDTrainer
```

```
from hydra.core.config_base import ConfigBase

@hydra.main(config_path="../configs", config_name="default", version_base=None)
def main(cfg: DictConfig):
    trainer = SGDTrainer(cfg)
    losses_gd, losses_mom = trainer.train()

    # Save artifacts (optional)
    import joblib, pathlib
    out = pathlib.Path("outputs")
    out.mkdir(exist_ok=True)
    joblib.dump({"gd": losses_gd, "mom": losses_mom}, out / "losses.pkl")

if __name__ == "__main__":
    main()
```

5. Thin Visualisation Notebook ([notebooks/demo.ipynb](#))

python ✖ 收起 自动换行

```
# Cell 1
%matplotlib inline
import joblib, matplotlib.pyplot as plt, numpy as np
data = joblib.load("../outputs/losses.pkl")

# Cell 2
maxiter = len(data["gd"])
plt.figure(figsize=(10,4))
```

```
plt.plot(np.log(np.abs(np.array(data["gd"])-data["gd"][-1])), 'r', label="GD")
plt.plot(np.log(np.abs(np.array(data["mom"])-data["mom"][-1])), 'b', label="momentum")
plt.title("Loss convergence (log scale)")
plt.xlabel("Iteration"); plt.ylabel("log(|loss_i - loss_opt|)")
plt.legend(); plt.grid(True); plt.show()
```

The notebook **never** contains training logic – only loads a pickle.

6. Testing ([tests/test_sgd.py](#))

pythonX 收起≡ 自动换行▷ 运行⌚ 复制

```
import torch
from src.optim.sgd import sgd_step
from src.models.logistic import LogisticRegression

def test_momentum_accumulates():
    model = LogisticRegression()
    x = torch.randn(10, 2)
    y = torch.randint(0, 2, (10,)).float()
    opt = lambda: torch.optim.SGD(model.parameters(), lr=0.1, momentum=0.9)

    # manual
    ...
```

```
velocity = 1j
for _ in range(2):
    model.zero_grad()
    loss = torch.nn.functional.binary_cross_entropy_with_logits(model(x), y)
    loss.backward()
    grads = [p.grad for p in model.parameters()]
    velocity = sgd_step(model.parameters(), grads, lr=0.1, momentum=0.9, velocity=)

# torch reference
ref = LogisticRegression()
ref_opt = opt()
for _ in range(2):
    ref.zero_grad()
    loss = torch.nn.functional.binary_cross_entropy_with_logits(ref(x), y)
    loss.backward()
    ref_opt.step()

assert all(torch.allclose(p1, p2) for p1, p2 in zip(model.parameters(), ref.parameters()))
```

Run with `pytest -q`.

7. Dependency Management

`requirements.txt`

```
text
```

torch>=2.2
hydra-core>=1.3
omegaconf>=2.3
wandb>=0.16
matplotlib
joblib
pytest

收起 自动换行 复制

`pyproject.toml (optional, for pip install -e .)`

```
toml
```

[project]
name = "q_sgd_momentum"
version = "0.1.0"
dependencies = [
 "torch>=2.2",

收起 自动换行 复制

```
"hydra-core>=1.3",
"omegaconf>=2.3",
"wandb>=0.16",
"matplotlib",
"joblib",
"pytest",
]
```

8. README (one-command demo)

markdown X 收起 ≡ 自动换行 ○ 复制

```
# SGD vs Momentum Demo

```bash
pip install -r requirements.txt
pip install -e .

GD only
python -m src.train optim.momentum=0.0

Momentum
python -m src.train optim.momentum=0.9

Open the plot
jupyter nbconvert --execute notebooks/demo.ipynb
```

All random seeds, hyper-parameters, and git hash are logged to **Weights & Biases** (disable with `logging.wandb.project=null`).

text X 收起 ≡ 自动换行 ○ 复制

---

```
9. Quick Checklist (paste into PR)
```

- [ ] Config externalised (no hard-coded numbers)
- [ ] set\_seed + deterministic flags
- [ ] Pure data generation function
- [ ] nn.Module model, versionable
- [ ] Optimiser is a pure function
- [ ] Trainer logs to wandb & console
- [ ] Losses saved as artifact (outputs/losses.pkl)
- [ ] Unit tests pass (pytest -q)
- [ ] black, isort, flake8, mypy clean
- [ ] README runs in < 5 seconds

text

✖ 收起

⇄ 自动换行

○ 复制

---

```
TL;DR
```

Original notebook	New engineering version
<hr/>	
All code in cells	**Package + Hydra config**
Hard-coded numbers	**External YAML**
No reproducibility	**Seed + deterministic**
No tests	**pytest suite**
Plotting mixed with training	**Thin notebook**
No logging	**wandb + console**

Copy the folder structure, `pip install -e .`, and you have a \*\*reproducible, testable,